

# CSE 163 $[\square, \circ, \Delta]$

## Numpy + Images

## Hunter Schafer




# What Happened?

```
a.shape = (3, 1)  
b.shape = (3, )
```

```
a.shape = (3, 1)  
b.shape = (1, 3)
```



```
a.shape = (3, 3)  
b.shape = (3, 3)
```

```
np.arange(3).reshape((3,1))+np.arange(3)
```



# Images as Matrices

Grey-scale images can be represented as matrices.



Grey-scale: 255




Grey-scale: 0

```
array([[ 84,   0,   0, ...,  31,  54,   0],
       [101,   0,   0, ...,   0,  32,   0],
       [  0,  89,   0, ...,  32,   0,   0],
       ...,
       [116,   0,   0, ...,   0,  83,   0],
       [  0,   0, 105, ...,  86,  97,  89],
       [103, 119, 124, ...,   0,   0,   0]])
```

# Color Images

When you overlap each color channel, it creates a picture we are used to seeing.



- Pixels on your monitor let out specified R/G/B light




# Color Images

Each pixel coordinate (  $x, y$  ) contains 3 values ranging for intensities of 0 to 255 ( 8-bit)  
- Red - Green - Blue

Mixing different intensities of each color gives us the full color spectrum.



5x4x3 NumPy Array



# Color Demo

**Goal:** Manipulate image values to change color of image

- Show how changing colors changes image
- Write a method to “circle” object in the middle of a picture

Note

- Often “global” transformations of an image can be done without loops, probably speeding up the program



Colab

## Brain Break



# Convolution

When wanting to use “local” information, we need to use a sliding window approach (i.e. a **convolution**)

Move the sliding window across the image, and compute the sum of the element wise product of the window (kernel) and image

Image

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

Kernel

0	1	2
2	2	0
0	1	2



# Convolution Example

3	3	2	1	0
0	0	1	3	1
3	1	$2_0$	$2_1$	$3_2$
2	0	$0_2$	$2_2$	$2_0$
2	0	$0_0$	$0_1$	$1_2$

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

# Common Kernels

What do the numbers in the kernel do?

Identity

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$



Edge Detection

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$




Sharpen

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Box Blur

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



Think 

1 minute

What is the result of applying a convolution using this kernel on this image?

Image

1	2	3
4	5	6
7	8	9

Kernel

1	1
0	2

[pollev.com/cse163](https://pollev.com/cse163)



What is the result of applying a convolution using this kernel on this image?

Image

1	2	3
4	5	6
7	8	9

Kernel

1	1
0	2




# Image Classification

For a really long time, image classification was done by painstakingly crafting these features (like edge detectors), by hand.

This kind of worked, but we quickly hit our peak using this method.


Then came the buzz-word... deep learning



# Convolutional Neural Networks

Has layers that perform functions like a convolution or classification

- Each layer is connected to the next (network)



Specifics are not important here, just general idea of learning kernels

# Image Classification

- Is this a solved problem?
  - We get pretty decent error rates on challenges like ImageNet
- What we can't do
  - Sometimes can't generalize to other real-world datasets
  - Adversarial attacks



"panda"  
57.7% confidence



“



**Eddy Dever**  
@EddyDever

Follow



It's terrifying that both of these things are true at the same time in this world:

- computers drive cars around
- the state of the art test to check that you're not a computer is whether you can successfully identify stop signs in pictures

12:26 AM - 13 May 2018

5,644 Retweets 12,727 Likes





# Why Convolutions?

A fundamental operation for image processing that shows up in many applications

- Edge detection
- Convolutional Neural Networks (CNN)
- Template Matching - Your Homework

